Generative AI Policy

Please note that this policy is a "work in progress" as the technology, the law and the Columbia community usage evolves.

PURPOSE

Columbia University is dedicated to advancing knowledge and learning, and embraces generative AI tools. The landscape of Generative AI is rapidly changing, and will change the way we teach, learn, and work. We encourage you to explore and experiment with these tools. The Office of the Provost has convened a working group of faculty and senior administrators from various parts of the University to develop policies and guidelines around the responsible use of these Generative AI tools (the "AI Team").

We ask that you review this guidance on the responsible use of generative AI in your work and study at Columbia University. Based on our collective experience with Generative AI use at the University, we anticipate that this guidance will evolve and be updated regularly.

Generative AI (or "AI") tools such as OpenAI's ChatGPT, Google's Bard, Stability AI's Stable Diffusion, and others, have captured the public's imagination as these tools become widely available for everyday use. Generative AI tools have the capacity to expedite existing processes and make possible new ones. These tools also have the potential to foster student learning and advance many aspects of research and health care delivery. While the University supports the responsible use of AI, these novel tools have notable limitations and present new risks that must be taken into consideration when using these technologies.

Two key attributes of these tools are the risk that an input could potentially become public, and the risk that the output may be biased, misleading, or inaccurate. There are risks related to information security, data privacy, copyright, and academic integrity and bias, for example:

- if Generative AI is given access to personal information, the technology may not respect the privacy rights of individuals, including in a manner that may be required for compliance with applicable data protection laws;
- if Generative AI is given access to confidential information or trade secrets, the University may lose its intellectual property (IP) rights to that information and the information may be disclosed to unauthorized

third parties through their independent use of the Generative AI technology;
- Generative AI outputs may violate the intellectual property rights of others, and might not themselves be protected by intellectual property laws;
- Generative AI outputs might be factually inaccurate, and we might be exposed to liability if we rely on those outputs without properly reviewing them; and
- Generative AI may produce decisions that are biased, discriminatory, or otherwise inconsistent with our policies, or that are otherwise in violation of applicable law.

In this initial policy, Columbia University requires that any use of Generative AI be in a manner reflective of its inherent limitations and to avoid these limitations and other emerging risks to the University, its faculty, researchers, students and staff and other stakeholders. Because AI is a rapidly evolving technology, the University will continue to monitor developments and will consider responses from the University community. This initial policy contains overarching guidelines that apply to all in the Columbia community while pursuing their Columbia activities. After these general requirements, the policy includes specific guidelines related to instruction (for both faculty and students) and research.

SCOPE

This Generative AI policy ("Policy") governs the use of Generative AI tools by staff, faculty, students, and researchers (the "Columbia community") in the performance of their functions for or on behalf of Columbia. Because this Policy may be updated from time to time, Columbia community members are encouraged to regularly review the most recent version of this Policy. Constructive comments from Columbia community members may be submitted here: AIpolicy@columbia.edu

DEFINITIONS

- **"Confidential Information"** means any business or technical information or research result belonging to Columbia, a Columbia community member, collaborators or other third parties, that is not publicly known or that has been provided or received under an obligation to maintain the information as confidential. Please note this

includes Protected Health Information or PHI. (see also the Information Security Charter)

- **"Generative AI"** includes any machine-based tool designed to consider user questions, prompts, and other inputs (e.g., text, images, videos) to generate a human-like output (e.g., a response to a question, a written document, software code, or a product design). Generative AI includes both standalone offerings such as ChatGPT, Bard, Stable Diffusion, and offerings that are embedded in other software, such as Github's Copilot.
- **"Personal Information"** means any information that, whether alone or in combination with other available information, identifies, relates to, describes, is reasonably capable of being associated with, or could reasonably be linked, directly or indirectly, to an individual. (see also definition of Personally Identifiable Information or PII in the Information Security Charter)

POLICY

Columbia expects all Columbia community members to follow these guidelines when using Generative AI tools for teaching and learning, research, and work-related functions:

- **Procuring AI Tools/Software (including free tools):** Contact Columbia University IT (CUIT) or CUIMC-IT, whichever is appropriate, before purchasing (or acquiring for free) AI products or products that contain functions that rely on AI to operate – especially when using University resources or University data (as defined in the Information Security Policy). CUIT's vendor management team will route the request to resources that can help validate the vendor's product and verify that the contract language does not introduce undue risk to the University. These processes can also direct you to existing vendors and potentially avoid duplicate spending. CUIT's vendor management team can be reached at cuitvmo@columbia.edu. CUIMC-IT's vendor management team can be reached through CUIT.
- **Do not input Confidential Information:** Columbia community members must not input any Confidential Information into Generative AI tools, except when permitted by validated contract language and security controls (approved by CUIT and central procurement).
- **Do not input Personal Information:** Columbia community members must not input any Personal Information about Columbia employees, students, faculty, or other stakeholders into a Generative

AI tool except when permitted by validated contract language and security controls (approved by CUIT and central procurement).

- **Do not input information that violates IP or general contract terms and conditions:** Columbia community members must be aware of the terms and conditions under which they are using AI tools. All members of the Columbia community must respect IP rights with the goal of protecting those IP rights. It is incumbent on the individual users to ensure that the inputs and the outputs of their AI tools are properly protected for reasons such as copyright and patent laws, data protection regulations, and identity theft crimes. Such AI tools might include Large Language Models (for example, ChatGPT), Machine Learning platforms (for example, AWS SageMaker), or image recognition software (for example, Google Cloud Vision). Please note that vendor licenses govern many of the digital resources provided by the Columbia University Libraries ("Libraries"), and some publishers are asserting that using their content with AI tools is not allowed. Please contact the Libraries for assistance in defining acceptable uses for licensed content with an AI tool or large language model here: ai-inquiries@library.columbia.edu.
- **Confirm the accuracy of the output provided by Generative AI tools:** Columbia community members must check the accuracy of information generated by Generative AI tools prior to relying on such information. Generative AI tools should not be relied upon without confirmation of accuracy from additional sources. It is possible for AI-generated content to be inaccurate, biased, or entirely fabricated (sometimes called "hallucinations"). Note that such AI-generated content may contain copyrighted material. You are responsible for any content that you publish that includes AI-generated material.
- **Check the output of Generative AI tools for bias:** Columbia community membersmust consider whether the data input into, and the output of, Generative AI tools produces decisions that may result in a disparate impact to individuals based on their protected classifications under applicable law, such as race, ethnicity, national origin, age, sexual orientation, or disability status. Do not rely on any output that is indicative of a potential bias.
- **Disclose the use of Generative AI tools:** Columbia community members who leverage Generative AI to produce any written materials or other work product must disclose that those materials and that work product is based on or derives from the use of Generative AI. Always be transparent if you are relying on the output of a Generative AI tool.
- **Comply with third-party intellectual property rights:** Columbia community membersmust not hold out any output generated by Generative AI tools as their own. If you quote, paraphrase or borrow ideas from the output of Generative AI tools, confirm that the output is

accurate and that you are not plagiarizing another party's existing work or otherwise violating another party's intellectual property rights.

- **Do not use Generative AI tools to produce malicious content:** Columbia community members are prohibited from using Generative AI tools to generate malicious content, such as malware, viruses, worms, and trojan horses that may have the ability to circumvent access control measures put in place by Columbia, or any other third-party entity, to prevent unauthorized access to their respective networks.
- **Ask the Generative AI system not to use inputs for training the system:** Some Generative AI systems permit users to opt out of the use of their data to train future iterations of the Generative AI system. Where that option is available, Columbia community members are expected to exercise it and opt out of training.

UNIVERSITY APPROVED APPLICATIONS, EXEMPTIONS, AND EXCEPTIONS

The AI Team will publish a list of specific applications that have been vetted by CUIT and Procurement and deemed fully or partially acceptable under this Policy. The list of these vetted Generative AI tools and the scope of the approval will be posted on the Office of the Provost website. Columbia community members are expected to use any approved Generative AI tools in accordance with the scope of this policy.

The AI Team may provide additional approval for AI tools where appropriate and for compelling business cases. Please contact AIpolicy@columbia.edu to raise requests for exceptions to this policy.

GENERATIVE AI AND ACADEMIC INTEGRITY

At Columbia University, it is our shared responsibility to promote intellectual honesty and scholarly integrity, which could be undermined with the utilization of AI-generated content being passed off as one's own work.

Particular considerations apply to use of AI by instructors and students. To ensure academic integrity, please refer to this guidance below.

For Students

The following guidance is shared to help support students in navigating the appropriate use of AI in their classes.

- Absent a clear statement from a course instructor granting permission, the use of Generative AI tools to complete an assignment or exam is prohibited. The unauthorized use of AI shall be treated similarly to unauthorized assistance and/or plagiarism (page 11 of Standards and Discipline).
- Students are encouraged to speak with their instructors regarding their expectations.

For Instructors

The following guidance is shared to help support faculty in navigating the appropriate use of AI in the classroom.

- At minimum, it is recommended that faculty share clear expectations at the beginning of each semester through the syllabus, policy distribution, and class discussion on the appropriate use of AI tools. Faculty can encourage students to reach out when they need support rather than risking a potential academic integrity violation. If permitted by the course, encourage students to acknowledge and cite any use of AI applications.
- It is recommended to develop a course policy about the use of AI tools and what faculty consider to be appropriate and inappropriate in their classes. For example, you may find it useful to include AI discussions in the classroom and online threads. Some instructors have partnered with students as they determine what constitutes appropriate use. This conversation and partnership can create opportunities for instructors and students to talk in detail about the evolution of particular tools, their potential benefits in specific disciplines, and their limitations. It is also an opportunity to be explicit about the course objectives and how the use of AI tools might interfere with or aid students' learning and their achievement of particular goals.
- Refer to introductory guidance and considerations for AI tools in the classroomfrom the Center for Teaching and Learning (CTL). Additionally, the CTL offers individual consultations, Faculty AI Labs, and AI Learning Communities by request from schools or departments.
- Important note about AI detection tools: Since the introduction of AI tools, there has been a parallel rise in tools claiming accurate detection of AI-generated work. As with any form of detection software, there are risks of misidentification, which can have consequences in the classroom. These products are best used with careful consideration and as one of many ways to work with students. It is also important to include the use of these tools in any discussion with students around

course policies, making clear why and how such services may be used in the course. As with other plagiarism detection tools, AI detection should be treated as a guideline and not a grading metric.

The University offers many support resources regarding academic integrity, for both instructors and students. You can find an overview of academic integrity resources–including considerations for AI tools–in the Promoting Academic Integrity resource, co-created and adapted from the faculty booklet Promoting Academic Integrity & Preventing Academic Dishonesty: Best Practices at Columbia University.

GENERATIVE AI AND RESEARCH

The policies above apply to all Columbia research activity. In addition, the following considerations and policies also extend to all research activity using AI:

- As with other tools and research methods, individuals who use Generative AI in research must be transparent regarding its use, in describing methods, acknowledgements, or elsewhere, as appropriate.
- Generative AI has been found to, among other things, generate citations to papers that do not exist by authors who do not exist, and has also been used to generate images for experiments that were never actually conducted. Researchers are responsible for the accuracy of any content created by AI that is included in any research output and must use caution in utilizing AI output in research.
- Researchers are expected to follow the policies of journals, funding agencies and professional societies through which they report their research. For example, some journals, such as Science, explicitly prohibit text, figures, images or graphics generated by ChatGPT or any other AI tools.
- Researchers must avoid uploading, or using as input, any unpublished research data or other Confidential Information into a Generative AI tool.
    - When a researcher inputs unpublished work of any kind into a Generative AI tool, the unpublished work becomes part of the universe of data in the AI. The model may incorporate the unpublished work into responses to queries from other researchers. Disclosure of the unpublished work to an AI tool may also impede or prevent future intellectual property protection for the unpublished work or give rise to privacy violations.
- Researchers must avoid uploading, or using as input, other parties' Confidential Information. Generative AI tools may not provide

protection for Confidential Information and their use could create the potential to breach confidential contractual commitments.
    - This includes, e.g., unpublished manuscripts or funding proposals that researchers may be asked to peer review. NIH and NSF, among others, prohibit using Generative AI for peer review.
    - This also includes the Personal Information of research subjects. For example, inputting interview data to perform preliminary analysis creates the possibility that quotations or other information from research subjects could become public, and potentially, that subjects could also be identified.
- Researchers should be mindful that the output of AI tools may infringe the rights of third parties since the responses generated are pooled from already established works.

Finally, all University research is subject to the University's research integrity policies, such as the Institutional Policy on Misconduct in Research and the Policy on Financial Conflicts of Interest and Research. More information about these policies is available on the Office of Research Compliance and Training website.

RELATED UNIVERSITY POLICIES

When using Generative AI tools, Columbia community members should always keep in mind the usual rules and policies concerning privacy, honor code, academic integrity, research conduct, information security and other such rules. Some of these applicable policies are as follows:

- Standards and Discipline
- University's Code of Academic Freedom and Tenure (page 53)
- Human Resources Policies
- Information Security Charter
- HIPAA
- Research Policies

RELATED UNIVERSITY RESOURCES

- Teaching and Learning in the Age of AI: Considerations, Resources, and Opportunities
- Incorporating Generative AI in Teaching: Faculty Examples Across Disciplines
- AI Community of Practice

QUESTIONS

If you have questions or concerns about this Policy or need guidance regarding your use of Generative AI tools please contact AIpolicy@columbia.edu.

APPENDIX

The Provost's Working Group on Generative AI is tasked to develop guidelines on the use (and procurement) of generative artificial intelligence tools (e.g. ChatGPT and Google Bard) by Columbia students, faculty, researchers and staff. More information about the Provost's Working Group, including a list of members, can be found here.